

Stage M2

Davide Ferre'

March 29, 2022

1 Useful stuff

1.1 Message Sequence Charts

Assume a finite set of processes \mathbb{P} and a finite set of messages \mathbb{M} . The set of (p2p) channels is $\mathbb{C} = \{(p, q) \in \mathbb{P} \times \mathbb{P} \mid p \neq q\}$. A send action is of the form $send(p, q, m)$ where $(p, q) \in \mathbb{C}$ and $m \in \mathbb{M}$. It is executed by p and sends message m to q . The corresponding receive action, executed by q , is $rec(p, q, m)$. For $(p, q) \in \mathbb{C}$, let $Send(p, q, _) = \{send(p, q, m) \mid m \in \mathbb{M}\}$ and $Rec(p, q, _) = \{rec(p, q, m) \mid m \in \mathbb{M}\}$. For $p \in \mathbb{P}$, we set $Send(p, _, _) = \{send(p, q, m) \mid q \in \mathbb{P} \setminus \{p\} \text{ and } m \in \mathbb{M}\}$, etc. Moreover, $\Sigma_p = Send(p, _, _) \cup Rec(_, p, _)$ will denote the set of all actions that are executed by p . Finally, $\Sigma = \bigcup_{p \in \mathbb{P}} \Sigma_p$ is the set of all the actions.

Peer-to-peer MSCs. A *p2p MSC* (or simply *MSC*) over \mathbb{P} and \mathbb{M} is a tuple $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ where \mathcal{E} is a finite (possibly empty) set of *events* and $\lambda : \mathcal{E} \rightarrow \Sigma$ is a labeling function. For $p \in \mathbb{P}$, let $\mathcal{E}_p = \{e \in \mathcal{E} \mid \lambda(e) \in \Sigma_p\}$ be the set of events that are executed by p . We require that \rightarrow (the *process relation*) is the disjoint union $\bigcup_{p \in \mathbb{P}} \rightarrow_p$ of relations $\rightarrow_p \subseteq \mathcal{E}_p \times \mathcal{E}_p$ such that \rightarrow_p is the direct successor relation of a total order on \mathcal{E}_p . For an event $e \in \mathcal{E}$, a set of actions $A \subseteq \Sigma$, and a relation $R \subseteq \mathcal{E} \times \mathcal{E}$, let $\#_A(R, e) = |\{f \in \mathcal{E} \mid (f, e) \in R \text{ and } \lambda(f) \in A\}|$. We require that $\triangleleft \subseteq \mathcal{E} \times \mathcal{E}$ (the *message relation*) satisfies the following:

- (1) for every pair $(e, f) \in \triangleleft$, there is a send action $send(p, q, m) \in \Sigma$ such that $\lambda(e) = send(p, q, m)$, $\lambda(f) = rec(p, q, m)$, and $\#_{Send(p, q, _)}(\rightarrow^+, e) = \#_{Rec(p, q, _)}(\rightarrow^+, f)$,
- (2) for all $f \in \mathcal{E}$ such that $\lambda(f)$ is a receive action, there is $e \in \mathcal{E}$ such that $e \triangleleft f$.

Finally, letting $\leq_M = (\rightarrow \cup \triangleleft)^*$, we require that \leq_M is a partial order. For convenience, we will simply write \leq when M is clear from the context.

Condition (1) above ensures that every (p2p) channel (p, q) behaves in a FIFO manner. By Condition (2), every receive event has a matching send event. Note that, however, there may be unmatched send events in an MSC. We let $SendEv(M) = \{e \in \mathcal{E} \mid \lambda(e) \text{ is a send action}\}$, $RecEv(M) = \{e \in \mathcal{E} \mid \lambda(e) \text{ is a receive action}\}$, $Matched(M) = \{e \in \mathcal{E} \mid \text{there is } f \in \mathcal{E} \text{ such that } e \triangleleft f\}$, and $Unm(M) = \{e \in \mathcal{E} \mid \lambda(e) \text{ is a send action and there is no } f \in \mathcal{E} \text{ such that } e \triangleleft f\}$. We do not distinguish isomorphic MSCs and let MSC_{p2p} be the set of all MSCs over the given sets \mathbb{P} and \mathbb{M} .

Example 1.1. For a set of processes $\mathbb{P} = \{p, q, r\}$ and a set of messages $\mathbb{M} = \{m_1, m_2, m_3, m_4\}$, $M_1 = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ is an MSC where, for example, $e_2 \triangleleft e'_2$ and $e'_3 \rightarrow e_4$. The dashed arrow means that the send event e_1 does not have a matching receive, so $e_1 \in Unm(M_1)$. Moreover, $e_2 \leq_{M_1} e_4$, but $e_1 \not\leq_{M_1} e_4$. We can find a total order $\rightsquigarrow \supseteq \leq_{M_1}$ such that $e_1 \rightsquigarrow e_2 \rightsquigarrow e'_2 \rightsquigarrow e_3 \rightsquigarrow e'_3 \rightsquigarrow e_4 \rightsquigarrow e'_4$. We call \rightsquigarrow a linearization, which is formally defined below.




Figure 1: MSC M_1

Mailbox MSCs. For an MSC $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$, we define an additional binary relation that represents a constraint under the mailbox semantics, where each process has only one incoming channel. Let $\sqsubset_M \subseteq \mathcal{E} \times \mathcal{E}$ be defined by: $e_1 \sqsubset_M e_2$ if there is $q \in \mathbb{P}$ such that $\lambda(e_1) \in Send(_, q, _)$, $\lambda(e_2) \in Send(_, q, _)$, and one of the following holds:

- $e_1 \in Matched(M)$ and $e_2 \in Unm(M)$, or
- $e_1 \triangleleft f_1$ and $e_2 \triangleleft f_2$ for some $f_1, f_2 \in \mathcal{E}_q$ such that $f_1 \rightarrow^+ f_2$.

We let $\preceq_M = (\rightarrow \cup \triangleleft \cup \sqsubset_M)^*$. Note that $\leq_M \subseteq \preceq_M$. We call $M \in \text{MSC}_{p2p}$ a *mailbox MSC* if \preceq_M is a partial order. Intuitively, this means that events can be scheduled in a way that corresponds to the mailbox semantics, i.e., with one incoming channel per process. Following the terminology in Bouajjani et al. 2018, we also say that a mailbox MSC satisfies *causal delivery*. The set of mailbox MSCs $M \in \text{MSC}_{p2p}$ is denoted by MSC_{mb} .

Example 1.2. MSC M_1 is a mailbox MSC. Indeed, even though the order \rightsquigarrow defined in Example 1.1 does not respect all mailbox constraints, particularly the fact that $e_4 \sqsubset_{M_1} e_1$, there is a total order $\rightsquigarrow \supseteq \preceq_{M_1}$ such that $e_2 \rightsquigarrow e_3 \rightsquigarrow e'_3 \rightsquigarrow e_4 \rightsquigarrow e_1 \rightsquigarrow e'_2 \rightsquigarrow e'_4$. We call \rightsquigarrow a mailbox linearization, which is formally defined below.

Linearizations, Prefixes, and Concatenation. Consider $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda) \in \text{MSC}$. A *p2p linearization* (or simply *linearization*) of M is a (reflexive) total order $\rightsquigarrow \subseteq \mathcal{E} \times \mathcal{E}$ such that $\leq_M \subseteq \rightsquigarrow$. Similarly, a *mailbox linearization* of M is a total order $\rightsquigarrow \subseteq \mathcal{E} \times \mathcal{E}$ such that $\preceq_M \subseteq \rightsquigarrow$. That is, every mailbox linearization is a p2p linearization, but the converse is not necessarily true (Example 1.2). Note that an MSC is a mailbox MSC iff it has at least one mailbox linearization.

Let $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda) \in \text{MSC}$ and consider $E \subseteq \mathcal{E}$ such that E is \leq_M -downward-closed, i.e, for all $(e, f) \in \leq_M$ such that $f \in E$, we also have $e \in E$. Then, the MSC $(E, \rightarrow \cap (E \times E), \triangleleft \cap (E \times E), \lambda')$, where λ' is the restriction of λ to E , is called a *prefix* of M . In particular, the empty MSC is a prefix of M . We denote the set of prefixes of M by $\text{Pref}(M)$. This is extended to sets $L \subseteq \text{MSC}$ as expected, letting $\text{Pref}(L) = \bigcup_{M \in L} \text{Pref}(M)$.

Lemma 1.1. *Every prefix of a mailbox MSC is a mailbox MSC.*

Proof. Let $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda) \in \text{MSC}_{mb}$ and $M_0 = (\mathcal{E}_0, \rightarrow_0, \triangleleft_0, \lambda_0)$ be a prefix of M , i.e., $\mathcal{E}_0 \subseteq \mathcal{E}$. By contradiction, suppose that M_0 is not a mailbox MSC. Then, there are distinct $e, f \in \mathcal{E}_0$ such that $e \preceq_{M_0} f \preceq_{M_0} e$ with $\preceq_{M_0} = (\rightarrow_0 \cup \triangleleft_0 \cup \sqsubset_{M_0})^*$. As $\mathcal{E}_0 \subseteq \mathcal{E}$, we have that $\rightarrow_0 \subseteq \rightarrow$, $\triangleleft_0 \subseteq \triangleleft$, and $\sqsubset_{M_0} \subseteq \sqsubset_M$. Finally, $\preceq_{M_0} \subseteq \preceq_M$ and M is not a mailbox MSC, which is a contradiction. \square

Let $M_1 = (\mathcal{E}_1, \rightarrow_1, \triangleleft_1, \lambda_1)$ and $M_2 = (\mathcal{E}_2, \rightarrow_2, \triangleleft_2, \lambda_2)$ be two MSCs. Their *concatenation* $M_1 \cdot M_2 = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ is defined if, for all $(p, q) \in \mathbb{C}$, $e_1 \in \text{Unm}(M_1)$, and $e_2 \in \mathcal{E}_2$ such that $\lambda(e_1) \in \text{Send}(p, q, _)$ and $\lambda(e_2) \in \text{Send}(p, q, _)$, we have $e_2 \in \text{Unm}(M_2)$. As expected, \mathcal{E} is the disjoint union of \mathcal{E}_1 and \mathcal{E}_2 , $\triangleleft = \triangleleft_1 \cup \triangleleft_2$, λ is the “union” of λ_1 and λ_2 , and $\rightarrow = \rightarrow_1 \cup \rightarrow_2 \cup R$. Here, R contains, for all $p \in \mathbb{P}$ such that $(\mathcal{E}_1)_p$ and $(\mathcal{E}_2)_p$ are non-empty, the pair (e_1, e_2) where e_1 is the maximal p -event in M_1 and e_2 is the minimal p -event in M_2 . Note that $M_1 \cdot M_2$ is indeed an MSC and that concatenation is associative.

1.2 Communicating Systems

We now recall the definition of communicating systems (aka communicating finite-state machines or message-passing automata), which consist of finite-state machines A_p (one for every process $p \in \mathbb{P}$) that can communicate through the FIFO channels from \mathbb{C} .

Definition 1.1. A *communicating system* over \mathbb{P} and \mathbb{M} is a tuple $\mathcal{S} = (A_p)_{p \in \mathbb{P}}$. For each $p \in \mathbb{P}$, $A_p = (\text{Loc}_p, \delta_p, \ell_p^0)$ is a finite transition system where Loc_p is a finite set of local (control) states, $\delta_p \subseteq \text{Loc}_p \times \Sigma_p \times \text{Loc}_p$ is the transition relation, and $\ell_p^0 \in \text{Loc}_p$ is the initial state.

Given $p \in \mathbb{P}$ and a transition $t = (\ell, a, \ell') \in \delta_p$, we let $\text{source}(t) = \ell$, $\text{target}(t) = \ell'$, $\text{action}(t) = a$, and $\text{msg}(t) = m$ if $a \in \text{Send}(_, _, m) \cup \text{Rec}(_, _, m)$.

There are in general two ways to define the semantics of a communicating system. Most often it is defined as a global infinite transition system that keeps track of the various local control states and all (unbounded) channel contents. As, in this paper, our arguments are based on a graph view of MSCs, we will define the language of \mathcal{S} directly as a set of MSCs. These two semantic views are essentially equivalent, but they have different advantages depending on the context. We refer to Aiswarya and Gastin 2014 for a thorough discussion.

Let $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ be an MSC. A *run* of \mathcal{S} on M is a mapping $\rho : \mathcal{E} \rightarrow \bigcup_{p \in \mathbb{P}} \delta_p$ that assigns to every event e the transition $\rho(e)$ that is executed at e . Thus, we require that (i) for all $e \in \mathcal{E}$, we have $\text{action}(\rho(e)) = \lambda(e)$, (ii) for all $(e, f) \in \rightarrow$, $\text{target}(\rho(e)) = \text{source}(\rho(f))$, (iii) for all $(e, f) \in \triangleleft$, $\text{msg}(\rho(e)) = \text{msg}(\rho(f))$, and (iv) for all $p \in \mathbb{P}$ and $e \in \mathcal{E}_p$ such that there is no $f \in \mathcal{E}$ with $f \rightarrow e$, we have $\text{source}(\rho(e)) = \ell_p^0$.

Letting run \mathcal{S} directly on MSCs is actually very convenient. This allows us to associate with \mathcal{S} its p2p language and mailbox language in one go. The *p2p language* of \mathcal{S} is $L_{p2p}(\mathcal{S}) = \{M \in \text{MSC}_{p2p} \mid$




Figure 2: System \mathcal{S}_1

there is a run of \mathcal{S} on M }. The *mailbox language* of \mathcal{S} is $L_{\text{mb}}(\mathcal{S}) = \{M \in \text{MSC}_{\text{mb}} \mid \text{there is a run of } \mathcal{S} \text{ on } M\}$.

Note that, following Bouajjani et al. 2018; Di Giusto, Laversa, and Lozes 2020, we do not consider final states or final configurations, as our purpose is to reason about all possible traces that can be *generated* by \mathcal{S} . The next lemma is obvious for the p2p semantics and follows from Lemma 1.1 for the mailbox semantics.

Lemma 1.2. *For all $\text{com} \in \{\text{p2p}, \text{mb}\}$, $L_{\text{com}}(\mathcal{S})$ is prefix-closed: $\text{Pref}(L_{\text{com}}(\mathcal{S})) \subseteq L_{\text{com}}(\mathcal{S})$.*

Example 1.3. Fig. 2 depicts $\mathcal{S}_1 = (A_p, A_q, A_r)$ such that MSC M_1 in Fig. 1 belongs to $L_{\text{p2p}}(\mathcal{S}_1)$ and to $L_{\text{mb}}(\mathcal{S}_1)$. There is a unique run ρ of \mathcal{S}_1 on M_1 . We can see that $(e'_3, e_4) \in \rightarrow$ and $\text{target}(\rho(e'_3)) = \text{source}(\rho(e_4)) = \ell_r^1$, $(e_2, e'_2) \in \triangleleft_{M_1}$, and $\text{msg}(\rho(e_2)) = \text{msg}(\rho(e'_2)) = m_2$.

1.3 Conflict Graph

We now recall the notion of a conflict graph associated to an MSC defined in Bouajjani et al. 2018. This graph is used to depict the causal dependencies between message exchanges. Intuitively, we have a dependency whenever two messages have a process in common. For instance, an \xrightarrow{SS} dependency between message exchanges v and v' expresses the fact that v' has been sent after v , by the same process. This notion is of interest because it was seen in Bouajjani et al. 2018 that the notion of synchronizability in MSCs (which is studied in this paper) can be graphically characterized by the nature of the associated conflict graph. It is defined in terms of linearizations in Di Giusto, Laversa, and Lozes 2020, but we equivalently express it directly in terms of MSCs.

For an MSC $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ and $e \in \mathcal{E}$, we define the type $\tau(e) \in \{S, R\}$ of e by $\tau(e) = S$ if $e \in \text{SendEv}(M)$ and $\tau(e) = R$ if $e \in \text{RecEv}(M)$. Moreover, for $e \in \text{Unm}(M)$, we let $\mu(e) = e$, and for $(e, e') \in \triangleleft$, we let $\mu(e) = \mu(e') = (e, e')$.

Definition 1.2 (Conflict graph). The *conflict graph* $\text{CG}(M)$ of an MSC $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ is the labeled graph $(\text{Nodes}, \text{Edges})$, with $\text{Edges} \subseteq \text{Nodes} \times \{S, R\}^2 \times \text{Nodes}$, defined by $\text{Nodes} = \triangleleft \cup \text{Unm}(M)$ and $\text{Edges} = \{(\mu(e), \tau(e)\tau(f), \mu(f)) \mid (e, f) \in \rightarrow^+\}$. In particular, a node of $\text{CG}(M)$ is either a single unmatched send event or a message pair $(e, e') \in \triangleleft$.

1.4 Logic and Special Tree-Width

Monadic Second-Order Logic. The set of MSO formulas over MSCs (over \mathbb{P} and \mathbb{M}) is given by the grammar $\varphi ::= x \rightarrow y \mid x \triangleleft y \mid \lambda(x) = a \mid x = y \mid x \in X \mid \exists x. \varphi \mid \exists X. \varphi \mid \varphi \vee \varphi' \mid \neg \varphi$, where $a \in \Sigma$, x and y are first-order variables, interpreted as events of an MSC, and X is a second-order variable, interpreted as a set of events. We assume that we have an infinite supply of variables, and we use common abbreviations such as \wedge, \forall , etc. The satisfaction relation is defined in the standard way and self-explanatory. For example, the formula $\neg \exists x. (\bigvee_{a \in \text{Send}(_, _, _)} \lambda(x) = a \wedge \neg \text{matched}(x))$ with $\text{matched}(x) = \exists y. x \triangleleft y$ says that there are no unmatched send events. It is not satisfied by MSC M_1 of Fig. 1, as message m_1 is not received, but by M_4 from Fig. ??.

Given a sentence φ , i.e., a formula without free variables, we let $L(\varphi)$ denote the set of (p2p) MSCs that satisfy φ . It is worth mentioning that the (reflexive) transitive closure of a binary relation defined by an MSO formula with free variables x and y , such as $x \rightarrow y$, is MSO-definable so that the logic can freely use formulas of the form $x \rightarrow^+ y$ or $x \leq y$ (where \leq is interpreted as \leq_M for the given MSC M). Therefore, the definition of a mailbox MSC can be readily translated into the formula $\varphi_{\text{mb}} = \neg \exists x. \exists y. (\neg(x = y) \wedge x \leq y \wedge y \leq x)$ so that we have $L(\varphi_{\text{mb}}) = \text{MSC}_{\text{mb}}$. Here, $x \leq y$ is obtained as the MSO-definable reflexive transitive closure of the union of the MSO-definable relations $\rightarrow, \triangleleft$, and \sqsubset . In particular, we may define $x \sqsubset y$ by :

$$x \sqsubset y = \bigvee_{\substack{q \in \mathbb{P} \\ a, b \in \text{Send}(_, q, _)}} \lambda(x) = a \wedge \lambda(y) = b \wedge \left(\begin{array}{l} \text{matched}(x) \wedge \neg \text{matched}(y) \\ \vee \exists x'. \exists y'. (x \triangleleft x' \wedge y \triangleleft y' \wedge x' \rightarrow^+ y') \end{array} \right)$$

Special Tree-Width. *Special tree-width* Courcelle 2010, is a graph measure that indicates how close a graph is to a tree (we may also use classical *tree-width* instead). This or similar measures are commonly employed in verification. For instance, tree-width and split-width have been used in Madhusudan and Parlato 2011 and, respectively, Cyriac, Gastin, and Kumar 2012; Aiswarya, Gastin, and Kumar 2014 to reason about graph behaviors generated by pushdown and queue systems. There are several ways to define the special tree-width of an MSC. We adopt the following game-based definition from Bollig and Gastin 2019.

Adam and Eve play a two-player turn based “decomposition game” whose positions are MSCs with some pebbles placed on some events. More precisely, Eve’s positions are *marked MSC fragments* (M, U) , where $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ is an *MSC fragment* (an MSC with possibly some edges from \triangleleft or \rightarrow removed) and $U \subseteq \mathcal{E}$ is the subset of marked events. Adam’s positions are pairs of marked MSC fragments. A move by Eve consists in the following steps:

1. marking some events of the MSC resulting in (M, U') with $U \subseteq U' \subseteq \mathcal{E}$,
2. removing (process and/or message) edges whose endpoints are marked,
3. dividing (M, U) in (M_1, U_1) and (M_2, U_2) such that M is the disjoint (unconnected) union of M_1 and M_2 and marked nodes are inherited.

When it is Adam’s turn, he simply chooses one of the two marked MSC fragments. The initial position is (M, \emptyset) where M is the (complete) MSC at hand. A terminal position is any position belonging to Eve such that all events are marked. For $k \in \mathbb{N}$, we say that the game is *k-winning* for Eve if she has a (positional) strategy that allows her, starting in the initial position and independently of Adam’s moves, to reach a terminal position such that, in every single position visited along the play, there are at most $k + 1$ marked events.

Fact 1.3 (Bollig and Gastin 2019). *The special tree-width of an MSC is the least k such that the associated game is k -winning for Eve.*

The set of MSCs whose special tree-width is at most k is denoted by $\text{MSC}^{k\text{-stw}}$.

1.5 Model Checking

In general, even simple verification problems, such as control-state reachability, are undecidable for communicating systems Brand and Zafiropulo 1983. However, they are decidable when we restrict to behaviors of bounded special tree-width, which motivates the following definition of a generic **bounded model-checking problem** for $\text{com} \in \{\text{p2p}, \text{mb}\}$:

Input: Two finite sets \mathbb{P} and \mathbb{M} , a communicating system \mathcal{S} , an MSO sentence φ , and $k \in \mathbb{N}$ (given in unary).

Question: Do we have $L_{\text{com}}(\mathcal{S}) \cap \text{MSC}^{k\text{-stw}} \subseteq L(\varphi)$?

Fact 1.4 (Bollig and Gastin 2019). *The bounded model-checking problem for $\text{com} = \text{p2p}$ is decidable. When the formulas φ are from LCPDL, then the problem is solvable in exponential time.*

Note that Bollig and Gastin 2019 does not employ the LCPDL modality jump, but it can be integrated easily. Using φ_{mb} or Φ_{mb} , we obtain the corresponding result for mailbox systems as a corollary:

Theorem 1.5. *The bounded model-checking problem for $\text{com} = \text{mb}$ is decidable. When the formulas φ are from LCPDL, then the problem is solvable in exponential time.*

1.6 Synchronizability

The above model-checking approach is incomplete in the sense that a positive answer does not imply correctness of the whole system. The system may still produce behaviors of special tree-width greater than k that violate the given property. However, if we know that a system only generates behaviors from a class whose special tree-width is bounded by k , we can still conclude that the system is correct.

This motivates the *synchronizability problem*. Several notions of synchronizability have been introduced in the literature. However, they all amount to asking whether all behaviors generated by a given communicating system have a particular shape, i.e., whether they are all included in a fixed (or given) set of MSCs \mathcal{C} . Thus, the synchronizability problem is essentially an inclusion problem, namely $L_{\text{p2p}}(\mathcal{S}) \subseteq \mathcal{C}$ or $L_{\text{mb}}(\mathcal{S}) \subseteq \mathcal{C}$. We show that, for decidability, it is enough to

Table 1: Summary of the decidability of the synchronizability problem in various classes

	PEER-TO-PEER	MAILBOX
Weakly synchronous	Undecidable [Thm. 1.10]	EXPTIME [Thm. 1.9]
Weakly k -synchronous	Decidable Bouajjani et al. 2018; Di Giusto, Laversa, and Lozes 2020 and [Thm. 1.10]	Decidable [Thm. 1.9]
Strongly k -synchronous	—	Decidable [Thm. ??]
Existentially k -p2p-bounded	Decidable Genest, Kuske, and Muscholl 2007, Prop. 5.5	Decidable [Thm. 1.9]
Existentially k -mailbox-bounded	—	Decidable [Prop. 1.6]

have that \mathcal{C} is MSO-definable and special-tree-width-bounded (STW-bounded): We call $\mathcal{C} \subseteq \text{MSC}$ (i) *MSO-definable* if there is an MSO-formula φ such that $L(\varphi) = \mathcal{C}$, (ii) *LCPDL-definable* if there is an LCPDL-formula Φ such that $L(\Phi) = \mathcal{C}$, (iii) *STW-bounded* if there is $k \in \mathbb{N}$ such that $\mathcal{C} \subseteq \text{MSC}^{k\text{-stw}}$.

An important component of the decidability proof is the following lemma, which shows that we can reduce synchronizability wrt. an STW-bounded class to bounded model-checking.

Lemma 1.6. *Let \mathcal{S} be a communicating system, $\text{com} \in \{\text{p2p}, \text{mb}\}$, $k \in \mathbb{N}$, and $\mathcal{C} \subseteq \text{MSC}^{k\text{-stw}}$. Then, $L_{\text{com}}(\mathcal{S}) \subseteq \mathcal{C}$ iff $L_{\text{com}}(\mathcal{S}) \cap \text{MSC}^{(k+2)\text{-stw}} \subseteq \mathcal{C}$.*

The result follows from the following lemma. Note that a similar property was shown in Genest, Kuske, and Muscholl 2007, Proposition 5.4 for the specific class of existentially k -bounded MSCs.

Lemma 1.7. *Let $k \in \mathbb{N}$ and $\mathcal{C} \subseteq \text{MSC}^{k\text{-stw}}$. For all $M \in \text{MSC} \setminus \mathcal{C}$, we have $(\text{Pref}(M) \cap \text{MSC}^{(k+2)\text{-stw}}) \setminus \mathcal{C} \neq \emptyset$.*

We now have all ingredients to state a generic decidability result for synchronizability:

Theorem 1.8. *Fix finite sets \mathbb{P} and \mathbb{M} . Suppose $\text{com} \in \{\text{p2p}, \text{mb}\}$ and let $\mathcal{C} \subseteq \text{MSC}$ be an MSO-definable and STW-bounded class (over \mathbb{P} and \mathbb{M}). The following problem is decidable: Given a communicating system \mathcal{S} , do we have $L_{\text{com}}(\mathcal{S}) \subseteq \mathcal{C}$?*

Proof. Consider the MSO-formula φ such that $L(\varphi) = \mathcal{C}$, and let $k \in \mathbb{N}$ such that $\mathcal{C} \subseteq \text{MSC}^{k\text{-stw}}$. We have $L_{\text{com}}(\mathcal{S}) \subseteq \mathcal{C} \xLeftrightarrow{\text{Lemma 1.6}} L_{\text{com}}(\mathcal{S}) \cap \text{MSC}^{(k+2)\text{-stw}} \subseteq \mathcal{C} \iff L_{\text{com}}(\mathcal{S}) \cap \text{MSC}^{(k+2)\text{-stw}} \subseteq L(\varphi)$. The latter can be solved thanks to Fact 1.4 and Theorem 1.5. \square

1.7 Application to Concrete Classes of Synchronizability

In this section, we instantiate our general framework by specific classes. Table 1 gives a summary of the results.

1.8 A New General Class: Weakly Synchronous MSCs

We first introduce the class of weakly synchronous MSCs. This is a generalization of synchronous MSCs studied earlier, in Bouajjani et al. 2018; Di Giusto, Laversa, and Lozes 2020, which we shall discuss later. We say an MSC is weakly synchronous if it is breakable into *exchanges* where an exchange is an MSC that allows one to schedule all sends before all receives. Let us define this formally:

Definition 1.3 (exchange). Let $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ be an MSC. We say that M is an *exchange* if $\text{SendEv}(M)$ is a \leq_M -downward-closed set.

Definition 1.4 (weakly synchronous). We say that $M \in \text{MSC}$ is *weakly synchronous* if it is of the form $M = M_1 \cdot \dots \cdot M_n$ such that every M_i is an exchange.

We use the term *weakly* to distinguish from variants introduced later.

Example 1.4. Consider the MSC M_2 in Fig. 3. It is weakly synchronous. Indeed, m_1 , m_2 , and m_5 are independent and can be put alone in an exchange. Repetitions of m_3 and m_4 are interlaced, but they constitute an exchange, as we can do all sends and then all receptions.

An easy adaptation of a characterization from Di Giusto, Laversa, and Lozes 2020 yields the following result for weakly synchronous MSCs:




Figure 3: MSC M_2

Proposition 1.1. Let M be an MSC. Then, M is weakly synchronous iff no RS edge occurs on any cyclic path in the conflict graph $\text{CG}(M)$.

It is easily seen that the characterization from Proposition 1.1 is LCPDL-definable:

Corollary 1.8.1. *The sets of weakly synchronous MSCs and weakly synchronous mailbox MSCs are LCPDL-definable. Both formulas have polynomial size.*

Moreover, under the mailbox semantics, we can show:

Proposition 1.2. The set of weakly synchronous mailbox MSCs is STW-bounded (in fact, it is included in $\text{MSC}^{4|\mathbb{P}|-\text{stw}}$).

Proof. Let M be fixed, and let us sketch Eve’s winning strategy. Let $n = |\mathbb{P}|$.

The first step for Eve is to split M in exchanges. She first disconnects the first exchange from the rest of the graph ($2n$ pebbles are needed), then she disconnects the second exchange from the rest of the graph ($2n$ pebbles needed, plus n pebbles remaining from the first round), and so on for each exchange.

So we are left with designing a winning strategy for Eve with $4n + 1$ pebbles on the graph of an exchange M_0 , where initially there are (at most) n pebbles placed on the first event of each process and also (at most) n pebbles placed on the last event of each process. Eve also places (at most) n pebbles on the last send event of each process and also (at most) n pebbles on the first receive event of each process. Eve erases the (at most) $n \rightarrow$ -edges between the last send event and the first receive event.

We are now in a configuration that will be our invariant.

Let us fix a mailbox linearization of M_0 and let e be the first send event in this linearization.

- if e is an unmatched send of process p , Eve places her last pebble on the next send event of p (if it exists), let us call it e' . Then Eve erases the \rightarrow -edge (e, e') , and now e is completely disconnected, so it can be removed and the pebble can be taken back.
- if $e \triangleleft e'$, with e' a receive event of process q , then due to the mailbox semantics e' is the first receive event of q , so it has a pebble placed on it. Eve removes the \triangleleft -edge between e and e' , then using the extra pebble she disconnects e and places a pebble on the \rightarrow -successor of e , then she also disconnects e' and places a pebble on the \rightarrow -successor of e' .

After that, we are back to our invariant, so we can repeat the same strategy with the second send event of the linearization, and so on until all edges have been erased. \square

We obtain the following result as a corollary. Note that it assumes the mailbox semantics.

Theorem 1.9. *The following problem is decidable in exponential time: Given \mathbb{P} , \mathbb{M} , and a communicating system \mathcal{S} (over \mathbb{P} and \mathbb{M}), is every MSC in $L_{\text{mb}}(\mathcal{S})$ weakly synchronous?*

Proof. According to Corollary 1.8.1, we determine the LCPDL formula Φ_{wsmb} such that $L(\Phi_{\text{wsmb}})$ is the set of weakly synchronous mailbox MSCs. Moreover, recall from Proposition 1.2 that the special tree-width of all weakly synchronous mailbox MSCs is bounded by $4|\mathbb{P}|$. By Lemma 1.6, $L_{\text{mb}}(\mathcal{S}) \subseteq L(\Phi_{\text{wsmb}})$ iff $L_{\text{mb}}(\mathcal{S}) \cap \text{MSC}^{(4|\mathbb{P}|+2)-\text{stw}} \subseteq L(\Phi_{\text{wsmb}})$. The latter is an instance of the bounded model-checking problem. As the length of Φ_{wsmb} is polynomial in $|\mathbb{P}|$, we obtain that the original problem is decidable in exponential time by Theorem 1.5. \square

For the same reasons, the model-checking problem for “weakly synchronous” systems is decidable. Interestingly, a reduction from Post’s correspondence problem shows that decidability fails when adopting the p2p semantics:

Theorem 1.10. *The following problem is undecidable: Given finite sets \mathbb{P} and \mathbb{M} as well as a communicating system \mathcal{S} , is every MSC in $L_{\text{p2p}}(\mathcal{S})$ weakly synchronous?*

Proof. We show that the control state reachability problem for p2p weakly synchronizable systems is not decidable. This immediately shows that the model-checking problem for p2p weakly synchronizable systems is not decidable.

D: Clarify why undecidability of control state reachability implies undecidability of model checking.

With some extra coding, it also shows that the membership problem (decide whether a given system is p2p weakly synchronizable) also is undecidable: indeed, it is enough to add a non weak

synchronizable behavior after the control states for which reachability is undecidable: the system will be not weakly synchronizable iff the control states are reached.

We reduce from Post correspondence problem (PCP). Let us recall that a PCP instance consists of N pairs (u_i, v_i) of finite words over an alphabet A , and that PCP undecidability holds already for $N = 7$ and $A = \{0, 1\}$. We let the set of messages be $\{1, \dots, N\} \uplus A \uplus \{\#\}$, and we consider a system with four machines: Prover1, Prover2, Verifier1, and Verifier2. We have unidirectional communication channels from provers to verifiers, so the system is weakly synchronous by construction.

Informally, the system works as follows:

- Prover1 guesses a solution $u_{i_1} \dots u_{i_m}$ of the PCP instance, and Prover2 also guesses the same solution $v_{i_1} \dots v_{i_m}$.
- Prover1 sends $u_{i_1} \dots u_{i_m}$ to Verifier1 and sends simultaneously $i_1 \dots i_m$ to Verifier2
- Prover2 sends $v_{i_1} \dots v_{i_m}$ to Verifier1 and sends simultaneously $i_1 \dots i_m$ to Verifier 2
- Verifier1 checks that the two words are equal and Verifier2 checks that the sequences of indices are equal.

Let us now formally define these machines. We describe them with regular expressions. For $w = a_1 \dots a_n$, we write $send^*(p, q, w)$ (resp $rec^*(p, q, w)$) for $send(p, q, a_1) \dots send(p, q, a_n)$ (resp $rec(p, q, a_1) \dots rec(p, q, a_n)$). We abbreviate Prover1 as P1, Prover2 as P2, Verifier1 as V1, and Verifier2 as V2

- Prover1 is

$$\left(\sum_{i=1}^N send(P_1, V_1, i) send^*(P_1, V_2, u_i) \right)^+ send(P_1, V_1, \#) send(P_1, V_2, \#)$$

- Prover2 is

$$\left(\sum_{i=1}^N send(P_2, V_1, i) send^*(P_2, V_2, v_i) \right)^+ send(P_2, V_1, \#) send(P_2, V_2, \#)$$

- Verifier1 is

$$\left(\sum_{i=1}^N rec(P_1, V_1, i) rec(P_2, V_1, i) \right)^* rec(P_1, V_1, \#) rec(P_2, V_1, \#)$$

- Verifier2 is

$$\left(\sum_{a \in \Sigma} rec(P_1, V_2, a) rec(P_2, V_2, a) \right)^* rec(P_1, V_2, \#) rec(P_2, V_2, \#)$$

It can be checked that all machines reach their own final state if and only if the PCP instance has a solution. \square

1.9 Weakly k -Synchronous MSCs

This negative result for the p2p semantics motivates the study of other classes. In fact, our framework captures several classes introduced in the literature.

Definition 1.5 (k -exchange). Let $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ be an MSC and $k \in \mathbb{N}$. We call M a k -exchange if M is an exchange and $|SendEv(M)| \leq k$.

Let us now recall the definition from Bouajjani et al. 2018; Di Giusto, Laversa, and Lozes 2020, but (equivalently) expressed directly in terms of MSCs rather than via *executions*. It differs from the weakly synchronous MSCs in that here, we insist on constraining the number of messages sent per exchange to be at most k .

Definition 1.6 (weakly k -synchronous). Let $k \in \mathbb{N}$. We say that $M \in \text{MSC}$ is weakly k -synchronous if it is of the form $M = M_1 \dots M_n$ such that every M_i is a k -exchange.

Example 1.5. MSC M_3 in Fig. 4 is weakly 1-synchronous, as it can be decomposed into three 1-exchanges (the decomposition is depicted by the horizontal dashed lines). We remark that $M_3 \in \text{MSC}_{\text{mb}}$. Note that there is a p2p linearization that respects the decomposition. On the other hand, a mailbox linearization needs to reorganize actions from different MSCs: the sending of m_3 needs to be done before the sending of m_1 . Note that M_1 in Fig. 1 is also weakly 1-synchronous.



Figure 4: MSC M_3

Proposition 1.3. Let $k \in \mathbb{N}$. The set of weakly k -synchronous p2p (mailbox, respectively) MSCs is effectively MSO-definable.

In fact, MSO-definability essentially follows from the following known theorem:

Theorem 1.11 (Di Giusto, Laversa, and Lozes 2020). *Let M be an MSC. Then, M is weakly k -synchronous iff every SCC in its conflict graph $\text{CG}(M)$ is of size at most k and no RS edge occurs on any cyclic path.*

This property is similar to the graphical characterization of weakly synchronous MSCs, except for the condition that every SCC in the conflict graph is of size at most k . Furthermore, it is easy to establish a bound on the special tree-width:

Proposition 1.4. Let $k \in \mathbb{N}$. The set of MSCs that are weakly k -synchronous have special tree-width bounded by $2k + |\mathbb{P}|$.

Hence, we can conclude that the class of weakly k -synchronous MSCs is MSO-definable and STW-bounded. As a corollary, we get the following (known) decidability result, but via an alternative proof:

Theorem 1.12 (Bouajjani et al. 2018; Di Giusto, Laversa, and Lozes 2020). *For $\text{com} \in \{\text{p2p}, \text{mb}\}$, the following problem is decidable: Given finite sets \mathbb{P} and \mathbb{M} , a communicating system \mathcal{S} , and $k \in \mathbb{N}$, is every MSC in $L_{\text{com}}(\mathcal{S})$ weakly k -synchronous?*

Proof. We proceed similarly to the proof of Theorem 1.9. For the given \mathbb{P} , \mathbb{M} , and k , we first determine, using Proposition 1.3, the MSO formula φ_k such that $L(\varphi_k)$ is the set of weakly k -synchronous p2p/mailbox MSCs. From Proposition 1.4, we know that the special tree-width of all weakly k -synchronous MSCs is bounded by $2k + |\mathbb{P}|$. By Lemma 1.6, we have $L_{\text{com}}(\mathcal{S}) \subseteq L(\varphi_k)$ iff $L_{\text{com}}(\mathcal{S}) \cap \text{MSC}^{(2k+|\mathbb{P}|+2)\text{-stw}} \subseteq L(\varphi_k)$. The latter is an instance of the bounded model-checking problem. By Fact 1.4 and Theorem 1.5, we obtain decidability. \square

Remark 1.1. The set of weakly k -synchronous MSCs is not directly expressible in LCPDL (the reason is that LCPDL does not have a built-in counting mechanism). However, its *complement* is expressible in the extension of LCPDL with existentially quantified propositions (we need $k + 1$ of them). The model-checking problem for this kind of property is still in EXPTIME and, therefore, so is the problem from Theorem 1.12 when k is given in unary. It is very likely that our approach can also be used to infer the PSPACE upper bound from Bouajjani et al. 2018 by showing bounded *path width* and using finite word automata instead of tree automata. Finally, note that the problem to decide whether there exists an integer $k \in \mathbb{N}$ such that all MSCs in $L_{\text{com}}(\mathcal{S})$ are weakly k -synchronous has recently been studied in Giusto, Laversa, and Lozes 2021 and requires different techniques.

Observe also that we can remove the constraint of all the sends preceding all the receives in a k -exchange, and still have decidability. We then have the following definition.

Definition 1.7 (modified k -exchange). Let $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ be an MSC and $k \in \mathbb{N}$. We call M a *modified k -exchange* if $|\text{SendEv}(M)| \leq k$.

We extend this notion to consider modified weakly k -synchronous executions as before, and the graphical characterization of this property is that there are at most k nodes in every SCC of the conflict graph. Hence, this class is also MSO-definable, and since each modified k -exchange has at most $2k$ events, it also has bounded special tree-width.

1.10 Existentially k -Bounded MSCs

Now, we turn to existentially k -bounded MSCs Lohrey and Muscholl 2002; Genest, Muscholl, and Kuske 2004; Genest, Kuske, and Muscholl 2007. Synchronizability has been studied for the p2p case in Genest, Kuske, and Muscholl 2007, so we only consider the mailbox case here. A linearization \rightsquigarrow of an MSC $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda) \in \text{MSC}$ is called *k -mailbox-bounded* if, for all $e \in \text{Matched}(M)$, say with $\lambda(e) = \text{send}(p, q, m)$, we have $\#_{\text{Send}(-, q, -)}(\rightsquigarrow, e) - \#_{\text{Rec}(-, q, -)}(\rightsquigarrow, e) \leq k$.

Definition 1.8. Let $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda) \in \text{MSC}$ and $k \in \mathbb{N}$. We call M *existentially k -mailbox-bounded* if it has some mailbox linearization that is k -mailbox-bounded.

Note that every existentially k -mailbox-bounded MSC is a mailbox MSC.

Example 1.6. MSC M_5 in Fig. 5 is existentially 1-mailbox-bounded, as witnessed by the (informally given) linearization $s(q, p, m_2) \rightsquigarrow s(p, q, m_1) \rightsquigarrow s(q, r, m_3) \rightsquigarrow r(q, r, m_3) \rightsquigarrow r(p, q, m_1) \rightsquigarrow s(p, q, m_1) \rightsquigarrow r(q, p, m_2) \rightsquigarrow s(q, r, m_3) \dots$. Note that M_5 is neither weakly nor strongly synchronous as we cannot divide it into exchanges.




Figure 5: MSC M_5

Proposition 1.5. For all $k \in \mathbb{N}$, the set of existentially k -p2p-bounded MSCs is MSO-definable and STW-bounded.

Proof. The set of existentially k -p2p-bounded MSCs was shown to be MSO-definable (in fact, even FO-definable) in Lohrey and Muscholl 2004. Note that there are minor differences in the definitions (in particular, the fact that we deal with unmatched messages), which, however, do not affect FO-definability. In Bollig and Gastin 2019, Proposition 5.4, page 163, it was shown that their special tree-width is bounded by $k|\mathbb{P}|^2 + |\mathbb{P}|$. \square

We obtain the following result as a corollary:

Theorem 1.13. For $\text{com} \in \{\text{p2p}, \text{mb}\}$, the following problem is decidable: Given finite sets \mathbb{P} and \mathbb{M} , a communicating system \mathcal{S} , and $k \in \mathbb{N}$, is every MSC in $L_{\text{com}}(\mathcal{S})$ existentially k -p2p-bounded?

Proof. Again, the proof follows exactly the same lines as that of Theorem 1.12, now using Proposition 1.5. \square

Note that this is similar to the problem considered in Genest, Kuske, and Muscholl 2007; Kuske and Muscholl 2014, though there is a subtle difference: in Genest, Kuske, and Muscholl 2007; Kuske and Muscholl 2014, there are a notion of deadlock and distinguished final configurations. We define the following relation in order to characterize k -mailbox-bounded MSCs.

Let $k \geq 1$, and let M be a fixed mailbox MSC. Let $\xrightarrow{\text{rev}}_k$ be the binary relation among events of M defined as follows: $r \xrightarrow{\text{rev}}_k s$ if

1. r is a receive event of a process p ;
2. let r' be the k -th receive event of process p after r ; then $s \triangleleft r'$.

Lemma 1.14. M is existential k -mailbox-bounded if and only if $\preceq_M \cup \xrightarrow{\text{rev}}_k$ is acyclic.

Proof. Assume that M is existential k -mailbox-bounded. Let \rightsquigarrow be a mailbox linearisation of M such that for all $e \in \text{Matched}(M)$, say with $\lambda(e) = \text{send}(p, q, m)$,

$$\#_{\text{Send}(-, q, -)}(\rightsquigarrow, e) - \#_{\text{Rec}(-, q, -)}(\rightsquigarrow, e) \leq k.$$

Then \rightsquigarrow is also a linearisation of $(\preceq_M \cup \xrightarrow{\text{rev}}_k)^*$. Indeed, if it was not the case, there would be a pair of events r, s such that $r \xrightarrow{\text{rev}}_k s$ and $s \rightsquigarrow r$. But then we would have

$$\#_{\text{Send}(-, q, -)}(\rightsquigarrow, s) - \#_{\text{Rec}(-, q, -)}(\rightsquigarrow, s) > k,$$

and the contradiction. So \rightsquigarrow is a linearisation of $(\preceq_M \cup \xrightarrow{\text{rev}}_k)^*$ and $\preceq_M \cup \xrightarrow{\text{rev}}_k$ is acyclic.

Conversely, assume that $\preceq_M \cup \xrightarrow{\text{rev}}_k$ is acyclic, and let \rightsquigarrow be a linearisation of $(\preceq_M \cup \xrightarrow{\text{rev}}_k)^*$. In particular, \rightsquigarrow is a mailbox linearisation of M . Let us show that for all $s \in \text{Matched}(M)$, say with $\lambda(s) = \text{send}(p, q, m)$,

$$\#_{\text{Send}(-, q, -)}(\rightsquigarrow, s) - \#_{\text{Rec}(-, q, -)}(\rightsquigarrow, s) \leq k.$$

Let $s \in \text{Matched}(M)$ be fixed, and let r' be such that $s \triangleleft r'$. There are two cases:

- $\#_{\text{Rec}(-,q,-)}(\rightarrow, r') \leq k$. Then

$$\#_{\text{Send}(-,q,-)}(\rightsquigarrow, s) \leq k,$$

because all sends before s are matched. So

$$\#_{\text{Send}(-,q,-)}(\rightsquigarrow, s) - \#_{\text{Rec}(-,q,-)}(\rightsquigarrow, s) \leq k,$$

- $\#_{\text{Rec}(-,q,-)}(\rightarrow, r') \leq k$. Then there is r on process q such that $r \xrightarrow{\text{rev}}_k s$. So $r \rightsquigarrow s$, and there are at most k messages in the buffer of q at the time of event s , or in other words,

$$\#_{\text{Send}(-,q,-)}(\rightsquigarrow, e) - \#_{\text{Rec}(-,q,-)}(\rightsquigarrow, e) \leq k.$$

So \rightsquigarrow is a mailbox linearisation with k bounded buffers, and M is existential k -mailbox-bounded. \square

Proposition 1.6. For all $k \in \mathbb{N}$, the set of existentially k -mailbox-bounded MSCs is MSO-definable and STW-bounded.

Proof. Let $k \geq 1$ be fixed. Since every existentially k -mailbox-bounded MSCs is also existentially k -p2p-bounded, and since the class of existentially k -p2p-bounded MSCs is STW bounded (cf Proposition 1.5), the class of existentially k -mailbox-bounded MSCs is also STW bounded.

Let us show that it is moreover MSO definable.

By Lemma 1.14, it is enough to show that the acyclicity of $\preceq_M \cup \xrightarrow{\text{rev}}_k$ is MSO definable, and since \preceq_M was already shown MSO definable and acyclicity is easily MSO definable, it is enough to show that $\xrightarrow{\text{rev}}_k$ is MSO definable. It is indeed the case, as demonstrated by this formula

$$\varphi(r, s) = \exists r_1, r_2, \dots, r_n. r \rightarrow r_1 \rightarrow r_2 \rightarrow \dots \rightarrow r_n \wedge s \triangleleft r_n.$$

Finally, let us show that existentially k -mailbox-bounded is also LCPDL definable. This follows from the following formulas:

$$\begin{aligned} \prec_M &= (\triangleleft + \rightarrow)^+ \\ R &= \langle \triangleleft^{-1} \rangle \top \\ \xrightarrow{\text{next } R} &= (\rightarrow \wedge \text{test}(\neg R))^* \cdot (\rightarrow \wedge \text{test}(R)) \\ \xrightarrow{\text{rev}}_k &= (\xrightarrow{\text{next } R})^k \cdot (\triangleleft)^{-1}. \\ \Phi_{\exists k \text{ mb-bounded}} &= \neg \text{ELoop} \langle (\prec_M + \xrightarrow{\text{rev}}_k)^+ \rangle \end{aligned}$$

\square

This extension is also valid for the p2p definition of existentially k -bounded MSCs, which were addressed in Genest, Kuske, and Muscholl 2007. Finally, our framework can also be adapted to treat universally bounded systems Henriksen et al. 2005; Lohrey and Muscholl 2002. Those extensions and the missing proofs are available in Appendix ??.

2 My stuff

2.1 Message Sequence Charts

Definition 2.1 (Causally ordered MSC). An MSC $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ is *causally ordered* if, for any two send events s and s' , such that $\lambda(s) = \text{Send}(_, q, _)$, $\lambda(s') = \text{Send}(_, q, _)$, and $s \leq_M s'$, we have either:

- $s, s' \in \text{Matched}(M)$ and $r \rightarrow^+ r'$, where r and r' are two receive events such that $s \triangleleft r$ and $s' \triangleleft r'$.
- $s' \in \text{Unm}(M)$.

2.2 Communicating Systems

Lemma 2.1. *Every prefix of a causally ordered MSC is a causally ordered MSC.*

Proof. Let $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda) \in \text{MSC}_{\text{co}}$ and let $M_0 = (\mathcal{E}_0, \rightarrow_0, \triangleleft_0, \lambda_0)$ be a prefix of M . By contradiction, suppose that M_0 is not a causally ordered MSC. There must be two distinct $s, s' \in \mathcal{E}_0$ such that $\lambda(s) = \text{Send}(_, q, _)$, $\lambda(s') = \text{Send}(_, q, _)$, $s \leq_{M_0} s'$ and either (i) $r' \rightarrow^+ r$, where r and r' are two receive events such that $s \triangleleft r$ and $s' \triangleleft r'$, or (ii) $s \in \text{Unm}(M_0)$ and $s' \in \text{Matched}(M_0)$. In both cases, M would also not be a causally ordered MSC, since $\mathcal{E}_0 \subseteq \mathcal{E}$, $\rightarrow_0 \subseteq \rightarrow$, and $\triangleleft_0 \subseteq \triangleleft$. This is a contradiction, thus M_0 has to be causally ordered. \square

Lemma 2.2 can be easily extended to $\text{com} = \text{co}$.

Lemma 2.2. *For all $\text{com} \in \{\text{p2p}, \text{mb}, \text{co}\}$, $L_{\text{com}}(\mathcal{S})$ is prefix-closed: $\text{Pref}(L_{\text{com}}(\mathcal{S})) \subseteq L_{\text{com}}(\mathcal{S})$.*

Proof. Follows from Lemma 2.1. \square

2.3 Model Checking

Proposition 2.1. The set MSC_{co} of causally ordered MSCs is MSO-definable.

Proof. Given an MSC M , it is causally ordered if it satisfies the MSO formula

$$\varphi_{\text{co}} = \neg \exists s. \exists s'. \left(\bigvee_{\substack{q \in \mathbb{P} \\ a, b \in \text{Send}(_, q, _)}} \lambda(s) = a \wedge \lambda(s') = b \wedge s \leq_M s' \wedge (\psi_1 \vee \psi_2) \right)$$

where ψ_1 and ψ_2 are

$$\psi_1 = \exists r. \exists r'. \left(\begin{array}{cc} s \triangleleft r & \wedge \\ s' \triangleleft r' & \wedge \\ r' \rightarrow^+ r & \end{array} \right) \quad \psi_2 = (\neg \text{matched}(s) \wedge \text{matched}(s'))$$

$$\text{matched}(x) = \exists y. x \triangleleft y$$

The property φ_{co} says that there cannot be two send events s and s' , with the same recipient, such that $s \leq_M s'$ and either (i) their corresponding receive events r and r' happen in the opposite order, i.e. $r' \rightarrow^+ r$, or (ii) s is unmatched and s' is matched. The set MSC_{co} of causally ordered MSCs is therefore MSO-definable as $\text{MSC}_{\text{co}} = L(\varphi_{\text{co}})$. \square

Knowing that MSC_{co} is MSO-definable, Theorem 1.5 can be restated for $\text{com} = \text{co}$.

Theorem 2.3. *The bounded model-checking problem for $\text{com} = \text{co}$ is decidable.*

Proof. By Proposition 2.1, $\text{MSC}_{\text{co}} = L(\varphi_{\text{co}})$. Given a system \mathcal{S} , we have that $L_{\text{co}}(\mathcal{S}) = L_{\text{p2p}}(\mathcal{S}) \cap L(\varphi_{\text{co}})$. Therefore, we can rewrite the bounded model checking problem for $\text{com} = \text{co}$ as

$$\begin{aligned} L_{\text{co}}(\mathcal{S}) \cap \text{MSC}^{k\text{-stw}} &\subseteq L(\varphi) \\ \iff L_{\text{p2p}}(\mathcal{S}) \cap L(\varphi_{\text{co}}) \cap \text{MSC}^{k\text{-stw}} &\subseteq L(\varphi) \\ \iff L_{\text{p2p}}(\mathcal{S}) \cap \text{MSC}^{k\text{-stw}} &\subseteq L(\varphi) \cup L(\neg \varphi_{\text{co}}) \\ \iff L_{\text{p2p}}(\mathcal{S}) \cap \text{MSC}^{k\text{-stw}} &\subseteq L(\varphi \vee \neg \varphi_{\text{co}}). \end{aligned}$$

The latter is decidable due to Fact 1.4. \square

2.4 Synchronizability

Note that Lemma 1.6 can be extended to $\text{com} = \text{co}$, since Lemma 1.7 does not depend on the kind of communication used by the system.

Lemma 2.4. *Let \mathcal{S} be a communicating system, $\text{com} \in \{\text{p2p}, \text{mb}, \text{co}\}$, $k \in \mathbb{N}$, and $\mathcal{C} \subseteq \text{MSC}^{k\text{-stw}}$. Then, $L_{\text{com}}(\mathcal{S}) \subseteq \mathcal{C}$ iff $L_{\text{com}}(\mathcal{S}) \cap \text{MSC}^{(k+2)\text{-stw}} \subseteq \mathcal{C}$.*

Theorem 1.8 can also be extended to $\text{com} = \text{co}$.

Theorem 2.5. *Fix finite sets \mathbb{P} and \mathbb{M} . Suppose $\text{com} \in \{\text{p2p}, \text{mb}, \text{co}\}$ and let $\mathcal{C} \subseteq \text{MSC}$ be an MSO-definable and STW-bounded class (over \mathbb{P} and \mathbb{M}). The following problem is decidable: Given a communicating system \mathcal{S} , do we have $L_{\text{com}}(\mathcal{S}) \subseteq \mathcal{C}$?*

Proof. Same as the proof for Theorem 1.8, but using Lemma 2.4 in place of Lemma 1.6, and Theorem 2.3 in place of Theorem 1.5. \square

Table 2: Summary of the decidability of the synchronizability problem in various classes

	P2P	CAUSAL ORDERING	MAILBOX
Weakly synchronous	Undecidable [Thm. 1.10]	Undecidable [Thm. 2.6]	EXPTIME [Thm. 1.9]
Weakly k -synchronous		Decidable [Thm. 2.7]	
Existentially k -bounded	Decidable	Decidable	Decidable

2.4.1 Weakly synchronous causally ordered MSCs

Corollary 1.8.1 can be extended to $\text{com} = \text{co}$.

Proposition 2.2. The set of weakly synchronous *causally ordered* MSCs is MSO-definable.

Proof. Both the sets of weakly synchronous MSCs and of causally ordered MSCs are MSO-definable, as shown by Corollary 1.8.1 and Proposition 2.1. Recall that any LCPDL-definable property is also MSO-definable. It suffices to take the conjunction of the two respective MSO formulas. \square

Theorem 2.6. *The following problem is undecidable: Given finite sets \mathbb{P} and \mathbb{M} as well as a communicating system \mathcal{S} , is every MSC in $L_{\text{co}}(\mathcal{S})$ weakly synchronous?*

Proof. The proof is essentially identical to the p2p case. We do the same reduction from the Post correspondence problem. Recall from the proof of Theorem 1.10 that we consider a system \mathcal{S} with four machines (P1, P2, V1, V2), where we have unidirectional communication channels from provers to verifiers. In particular notice that all the possible behaviours of \mathcal{S} are causally ordered, i.e. $L_{\text{p2p}}(\mathcal{S}) \subseteq \text{MSC}_{\text{co}}$; according to how we built our system \mathcal{S} , it is impossible to have a pair of causally-related send events of P1 and P2¹, which implies that causal ordering is already ensured by any possible p2p behaviour of \mathcal{S} . The rest of the proof is identical to the p2p case. \square

Corollary 2.6.1. *The set of weakly synchronous causally ordered MSCs has unbounded special tree-width.*

Proof. Suppose that the set of weakly synchronous causally ordered MSCs is STW-bounded. By Proposition 2.2 and Theorem 2.5, we have that the synchronicity problem for the class of weakly synchronous causally ordered MSCs would be decidable. This is a contradiction, since Theorem 2.6 states that this problem is undecidable. \square

2.4.2 Weakly k -synchronous causally ordered MSCs

Proposition 2.3. The set of weakly k -synchronous causally ordered MSCs is MSO-definable.

Proof. Both the sets of weakly k -synchronous MSCs and of causally ordered MSCs are MSO-definable, as shown by Proposition 1.3 and Proposition 2.1. It suffices to take the conjunction of the two respective MSO formulas. \square

Theorem 1.8 can be easily extended to $\text{com} = \text{co}$.

Theorem 2.7. *For $\text{com} \in \{\text{p2p}, \text{mb}, \text{co}\}$, the following problem is decidable: Given finite sets \mathbb{P} and \mathbb{M} , a communicating system \mathcal{S} , and $k \in \mathbb{N}$, is every MSC in $L_{\text{com}}(\mathcal{S})$ weakly k -synchronous?*

Proof. By Proposition 2.3 and Proposition 1.4 we have that the class of causally ordered k -synchronous MSCs is MSO-definable and STW-bounded². Theorem 2.5 ends the proof. \square

2.4.3 Existentially k causally ordered bounded MSCs

Definition 2.2. Let $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda) \in \text{MSC}$ and $k \in \mathbb{N}$. A linearization \rightsquigarrow of M is called *k -bounded* if, for all $e \in \text{Matched}(M)$, with $\lambda(e) = \text{send}(p, q, m)$, we have

$$\#_{\text{Send}(p, q, _)}(\rightsquigarrow, e) - \#_{\text{Rec}(p, q, _)}(\rightsquigarrow, e) \leq k.$$

Recall that $\#_{\text{Send}(p, q, _)}(\rightsquigarrow, e)$ denotes the number of send events from p to q that occurred before e , according to \rightsquigarrow .

¹There is no channel between P1 and P2, and we only have unidirectional communication channels from provers to verifiers; it is impossible to have a causal path between two send events of P1 and P2.

²Note that Proposition 1.4 is independent from the type of communication.

Definition 2.3. An MSC is said to be *existentially $p2p$ bounded* ($\exists k\text{-}p2p\text{-bounded}$) if it has a k -bounded linearization.

Definition 2.4. An MSC is said to be *existentially k causally ordered bounded* ($\exists k\text{-co-bounded}$) if it is causally ordered and it has a k -bounded linearization.

Note that every existentially k causally ordered bounded MSC is an existentially $k\text{-}p2p\text{-bounded}$ MSC.

Proposition 2.4. For all $k \in \mathbb{N}$, the set of $\exists k\text{-co-bounded}$ MSCs is MSO-definable and STW-bounded.

Proof. Let $\text{MSC}_{\exists k\text{-}p2p\text{-}b}$ and $\text{MSC}_{\exists k\text{-co-b}}$ be the set of existentially $k\text{-}p2p\text{-bounded}$ MSCs and the set of existentially k causally ordered bounded MSCs, respectively. $\text{MSC}_{\exists k\text{-}p2p\text{-}b}$ was shown to be both MSO-definable (in Lohrey and Muscholl 2004) and STW-bounded (in Bollig and Gastin 2019, Proposition 5.4, page 163). $\text{MSC}_{\exists k\text{-co-b}}$ also has to be STW-bounded, since we have $\text{MSC}_{\exists k\text{-co-b}} \subseteq \text{MSC}_{\exists k\text{-}p2p\text{-}b}$. Note that, by definition, $\text{MSC}_{\exists k\text{-co-b}} = \text{MSC}_{\exists k\text{-}p2p\text{-}b} \cap \text{MSC}_{\text{co}}$. Since both $\text{MSC}_{\exists k\text{-}p2p\text{-}b}$ and MSC_{co} can be defined by an MSO formula, the latter according to Proposition 2.1, $\text{MSC}_{\exists k\text{-co-b}}$ is also MSO-definable³. \square

Theorem 2.8. *The following problem is decidable: Given finite sets \mathbb{P} and \mathbb{M} , a communicating system \mathcal{S} , and $k \in \mathbb{N}$, is every MSC in $L_{\text{co}}(\mathcal{S})$ $\exists k\text{-co-bounded}$?*

Proof. Directly follows from 2.4 and 2.5. \square

3 Sketches

3.1 MSCs and partial order

D: Give intuition of what is an MSC and how MSCs can be used to represent graphically the behaviour of a system in terms of send/receive events. Talk about partial order on MSCs (causal relation/paths) and linearizations.

D: Clarify difference between 'message delivery' and 'receive event'. (Actually there should not be any difference according to our interpretation)

3.2 Communication architectures and variants

A communicating system, intended as a set of finite-state machines that can exchange messages through channels, may use different *communication architectures*. We will consider the following:

- *Fully asynchronous*: a fully asynchronous architecture (or simply asynchronous from now on) can be modeled as a collection of channels between each pair (M_1, M_2) of machines, such that M_1 can send messages to M_2 . Following this description, given two machines M_1 and M_2 that can exchange messages in both directions, we have two channels: one for the messages sent by M_1 to M_2 , and the other for the messages sent by M_2 to M_1 . The channels behave as *bags*, which means that they do not guarantee any specific order on the delivery of messages. This architecture is equivalent to the "Fully asynchronous" communication model in Chevrou, Hurault, and Quéinnec 2016.

D: Is it true??

- *FIFO 1–1 ($p2p$)*: this is a variant of the asynchronous architecture in which channels operate in FIFO mode (i.e. as queues). This means that messages between a couple of peers are delivered in their send order, whereas messages from/to different peers are delivered independently. We will use the term *peer-to-peer* ($p2p$) as a synonym of FIFO 1–1. This architecture is equivalent to the "FIFO 1–1" communication model in Chevrou, Hurault, and Quéinnec 2016.

³Suppose $\varphi_{\exists k\text{-}p2p\text{-}b}$ is the MSO formula for $\text{MSC}_{\exists k\text{-}p2p\text{-}b}$, and φ_{co} is the MSO formula for MSC_{co} . Then, $\text{MSC}_{\exists k\text{-co-b}}$ is defined by $\varphi_{\exists k\text{-co-b}} = \varphi_{\exists k\text{-}p2p\text{-}b} \wedge \varphi_{\text{co}}$

- *Causally ordered*: this can be described as a more specific variant of the p2p architecture. In a causally ordered architecture, the communicating system ensures that messages are delivered according to the causality of their emissions. In other words, if a message m_1 is causally sent before a message m_2 (i.e. there exists a causal path from the first emission to the second one), then a peer cannot receive m_2 before m_1 . Channels still operate in FIFO mode, but the delivery of some messages might be delayed to enforce causal ordering. This architecture is equivalent to the "Causally ordered" communication model in Chevrout, Hurault, and Quéinnec 2016. In literature, several implementations of causal ordering have been proposed. For instance, the algorithm described in Schiper, Eggli, and Sandoz 1989 makes use of the logical vector clocks introduced by Mattern-Fidge Fidge 1988; Mattern 1989 to enforce causal ordering.
- *FIFO n-1 (mailbox)*: in a mailbox architecture, each machine merges all of its incoming messages (from any source) into a unique queue. In other words, we can model it as each machine P_i having a single incoming FIFO channel, which is shared by the other machines that can send messages to P_i . A send event consists in adding the message at the end of the queue of the destination peer. This architecture is equivalent to the "FIFO n-1" communication model in Chevrout, Hurault, and Quéinnec 2016.
- *FIFO n-n*: this architecture can be modeled as a unique shared FIFO channel. Messages are globally ordered and delivered according to their emission order. As noted in Chevrout, Hurault, and Quéinnec 2016, this model is generally unrealistic and its implementations are inefficient.

For convenience, we will refer to a system that uses the p2p architecture simply as a p2p system. The same shorthand will be used for the other communication architectures. Note that, for each of these communicating architectures, there may be other equivalent ways of describing/modeling them.

3.3 Definitions

A Message Sequence Chart (MSC), such as the one in Fig. ??, provides a visual description of the behaviour of a distributed system. In this section, we start by formally defining the most generic class of Message Sequence Charts (MSCs), which we call asynchronous MSCs. More specialized classes of MSCs, such as *p2p* MSCs, will also be discussed. Intuitively, we say that an MSC M is asynchronous if there is an asynchronous system \mathcal{S} that can exhibit the behaviour described by M .

Definition 3.1 (Asynchronous MSC). An *asynchronous MSC* over \mathbb{P} and \mathbb{M} is a tuple $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$, where \mathcal{E} is a finite (possibly empty) set of *events* and $\lambda : \mathcal{E} \rightarrow \Sigma$ is a labeling function that associates an action to each event. For $p \in \mathbb{P}$, let $\mathcal{E}_p = \{e \in \mathcal{E} \mid \lambda(e) \in \Sigma_p\}$ be the set of events that are executed by p . We require that \rightarrow (the *process relation*) is the disjoint union $\bigcup_{p \in \mathbb{P}} \rightarrow_p$ of relations $\rightarrow_p \subseteq \mathcal{E}_p \times \mathcal{E}_p$ such that \rightarrow_p is the direct successor relation of a total order on \mathcal{E}_p . For an event $e \in \mathcal{E}$, a set of actions $A \subseteq \Sigma$, and a relation $R \subseteq \mathcal{E} \times \mathcal{E}$, let $\#_A(R, e) = |\{f \in \mathcal{E} \mid (f, e) \in R \text{ and } \lambda(f) \in A\}|$. We require that $\triangleleft \subseteq \mathcal{E} \times \mathcal{E}$ (the *message relation*) satisfies the following:

- (1) for every pair $(e, f) \in \triangleleft$, there is a send action $send(p, q, m) \in \Sigma$ such that $\lambda(e) = send(p, q, m)$, $\lambda(f) = rec(p, q, m)$.
- (2) for all $f \in \mathcal{E}$ such that $\lambda(f)$ is a receive action, there is $e \in \mathcal{E}$ such that $e \triangleleft f$.

Finally, letting $\leq_M = (\rightarrow \cup \triangleleft)^*$, we require that \leq_M is a partial order. For convenience, we simply write \leq when M is clear from the context. We will refer to \leq as the *causal ordering* or *happens-before* relation.

According to Condition (2), every receive event must have a matching send event. Note that, however, there may be unmatched send events. We let $SendEv(M) = \{e \in \mathcal{E} \mid \lambda(e) \text{ is a send action}\}$, $RecEv(M) = \{e \in \mathcal{E} \mid \lambda(e) \text{ is a receive action}\}$, $Matched(M) = \{e \in \mathcal{E} \mid \text{there is } f \in \mathcal{E} \text{ such that } e \triangleleft f\}$, and $Unm(M) = \{e \in \mathcal{E} \mid \lambda(e) \text{ is a send action and there is no } f \in \mathcal{E} \text{ such that } e \triangleleft f\}$. We do not distinguish isomorphic MSCs and let MSC_{as} be the set of all the asynchronous MSCs over the given sets \mathbb{P} and \mathbb{M} .

Linearizations. Consider $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda) \in \text{MSC}_{\text{as}}$. A *linearization* of M is a (reflexive) total order $\rightsquigarrow \subseteq \mathcal{E} \times \mathcal{E}$ such that $\leq_M \subseteq \rightsquigarrow$. In other words, a linearization of M is a total order that respects the happens-before relation \leq_M defined over M .

D: Provide example of linearization.

Asynchronous MSCs are the widest class of MSCs that we will deal with. By introducing additional constraints, we are able to define other classes of MSCs that exclusively describe the behaviours of p2p systems, mailbox systems, and so on.

D: Provide example of asynchronous MSC that is not p2p.

As in the asynchronous case, we say that M is a p2p MSC if there is a p2p system that can produce the behaviour described by M . We give here the formal definition of p2p MSC, which also considers the possibility of having unmatched messages (i.e. messages that are sent but not received).

D: Why unmatched messages? What do they represent?

Definition 3.2 (Peer-to-peer MSCs). A *p2p MSC* (or simply *MSC*) is an asynchronous MSC where we require that, for every pair $(e, f) \in \triangleleft$, such that $\lambda(e) = \text{send}(p, q, m)$, $\lambda(f) = \text{rec}(p, q, m)$, we have $\#_{\text{Send}(p, q, _)}(\rightarrow^+, e) = \#_{\text{Rec}(p, q, _)}(\rightarrow^+, f)$.

The additional constraint satisfied by p2p MSCs ensures that channels operate in FIFO mode; when a process q receives a message from a process p , it must have already received all the messages that were previously sent to him by p . Let MSC_{p2p} denote the set of all the p2p MSCs over two given sets \mathbb{P} and \mathbb{M} . Note that, by definition, every (p2p) MSC is an asynchronous MSC. The idea is that we are always able to find an asynchronous system that *can* exhibit the behaviour described by an MSC; after all, the channels of an asynchronous system do not have to follow any specific behaviour, so they can indeed happen to operate as if they were queues. Example ?? shows that the opposite direction is generally not true, an asynchronous MSC is not always an MSC. It follows that $\text{MSC}_{\text{p2p}} \subset \text{MSC}_{\text{as}}$.

We will now consider the class of MSCs for which there is a causally ordered system that can produce their behaviour. Intuitively, an MSC is causally ordered if all the messages sent to the same process are received in an order which is consistent with the causal ordering of the corresponding send events. Below the formal definition, which also considers unmatched messages.

D: Provide example of an MSC that is not causally ordered

Definition 3.3 (Causally ordered MSC). An MSC $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ is *causally ordered* if, for any two send events s and s' , such that $\lambda(s) = \text{Send}(_, q, _)$, $\lambda(s') = \text{Send}(_, q, _)$, and $s \leq_M s'$, we have either:

- $s, s' \in \text{Matched}(M)$ and $r \rightarrow^+ r'$, where r and r' are two receive events such that $s \triangleleft r$ and $s' \triangleleft r'$.
- $s' \in \text{Unm}(M)$.

By definition, every causally ordered MSC is a (p2p) MSC. This is not surprising, considering that a causally ordered system is essentially a p2p system with an additional constraint on the delivery of messages; indeed, we are always able to find a p2p system that *can* exhibit the behaviour described by a causally ordered MSC. Let MSC_{co} denote the set of all the causally ordered MSCs over two given sets \mathbb{P} and \mathbb{M} . Example ?? shows that an MSC is not always a causally ordered MSC. It follows that $\text{MSC}_{\text{co}} \subset \text{MSC}_{\text{p2p}}$.

Moving on to the mailbox semantics, we say that M is a mailbox MSC if there is a mailbox system that can exhibit the behaviour described by M .

D: Provide example of MSC which is not mailbox

Definition 3.4 (Mailbox MSC). An MSC $M = (\mathcal{E}, \rightarrow, \triangleleft, \lambda)$ is a *mailbox MSC* if it has a linearization \rightsquigarrow where, for any two send events s and s' , such that $\lambda(s) = \text{Send}(_, q, _)$, $\lambda(s') = \text{Send}(_, q, _)$, and $s \rightsquigarrow s'$, we have either:

- $s, s' \in \text{Matched}(M)$ and $r \rightsquigarrow r'$, where r and r' are two receive events such that $s \triangleleft r$ and $s' \triangleleft r'$.
- $s' \in \text{Unm}(M)$.

Such a linearization will be referred to as a *mailbox linearization*. Let MSC_{mb} denote the set of all the mailbox MSCs over two given sets \mathbb{P} and \mathbb{M} . By definition, every mailbox MSC is a (p2p) MSC. Conversely, Example ?? shows an MSC which is not a mailbox MSC. It follows that $\text{MSC}_{\text{mb}} \subset \text{MSC}_{\text{p2p}}$. We show here that each mailbox MSC is also a causally ordered MSC.

Proposition 3.1. Every mailbox MSC is a causally ordered MSC.

Proof. Let M be a mailbox MSC and \rightsquigarrow a mailbox linearization of it. Recall that a linearization has to respect the happens-before partial order over M , i.e. $\leq_M \subseteq \rightsquigarrow$. Consider any two send events s and s' , such that $\lambda(s) = \text{Send}(_, q, _)$, $\lambda(s') = \text{Send}(_, q, _)$ and $s \leq_M s'$. Since $\leq_M \subseteq \rightsquigarrow$, we have that $s \rightsquigarrow s'$ and, by the definition of mailbox linearization, either (i) $s' \in \text{Unm}(M)$, or (ii) $s, s' \in \text{Matched}(M)$, $s \triangleleft r$, $s' \triangleleft r'$ and $r \rightsquigarrow r'$. The former clearly respects the definition of causally ordered MSC, so let us focus on the latter. Note that r and r' are two receive events executed by the same process, hence $r \rightsquigarrow r'$ implies $r \rightarrow^+ r'$. It follows that M is a causally ordered MSC. \square

D: Provide example of causally ordered MSC which is not mailbox

From Proposition 3.1 and Example ?? it follows that $\text{MSC}_{\text{mb}} \subset \text{MSC}_{\text{co}}$. The four classes of MSCs that we presented form a hierarchy, namely $\text{MSC}_{\text{mb}} \subset \text{MSC}_{\text{co}} \subset \text{MSC}_{\text{p2p}} \subset \text{MSC}_{\text{as}}$, as shown by Fig. 6.

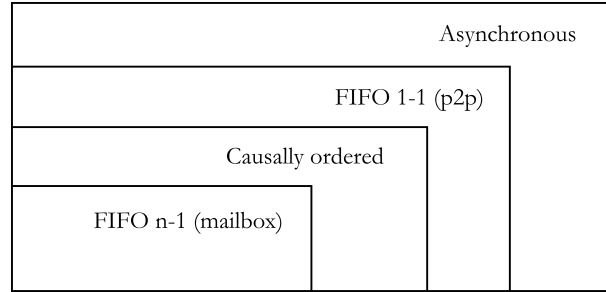


Figure 6: The hierarchy of MSC classes.

References

- Aiswarya, C. and Paul Gastin (2014). “Reasoning About Distributed Systems: WYSIWYG (Invited Talk)”. In: *34th International Conference on Foundation of Software Technology and Theoretical Computer Science, FSTTCS 2014, December 15-17, 2014, New Delhi, India*. Ed. by Venkatesh Raman and S. P. Suresh. Vol. 29. LIPIcs. Schloss Dagstuhl - Leibniz-Zentrum für Informatik, pp. 11–30. DOI: [10.4230/LIPIcs.FSTTCS.2014.11](https://doi.org/10.4230/LIPIcs.FSTTCS.2014.11). URL: <https://doi.org/10.4230/LIPIcs.FSTTCS.2014.11>.
- Aiswarya, C., Paul Gastin, and K. Narayan Kumar (2014). “Verifying Communicating Multipushdown Systems via Split-Width”. In: *Automated Technology for Verification and Analysis - 12th International Symposium, ATVA 2014*. Vol. 8837. Lecture Notes in Computer Science. Springer, pp. 1–17.
- Bollig, Benedikt and Paul Gastin (2019). “Non-Sequential Theory of Distributed Systems”. In: *CoRR* abs/1904.06942. arXiv: [1904.06942](https://arxiv.org/abs/1904.06942). URL: <http://arxiv.org/abs/1904.06942>.
- Bouajjani, Ahmed et al. (2018). “On the Completeness of Verifying Message Passing Programs Under Bounded Asynchrony”. In: *Computer Aided Verification - 30th International Conference, CAV 2018, Held as Part of the Federated Logic Conference, FloC 2018, Oxford, UK, July 14-17, 2018, Proceedings, Part II*. Ed. by Hana Chockler and Georg Weissenbacher. Vol. 10982. Lecture Notes in Computer Science. Springer, pp. 372–391. DOI: [10.1007/978-3-319-96142-2_23](https://doi.org/10.1007/978-3-319-96142-2_23). URL: https://doi.org/10.1007/978-3-319-96142-2_23.
- Brand, Daniel and Pitro Zafiropulo (1983). “On Communicating Finite-State Machines”. In: *J. ACM* 30.2, pp. 323–342. DOI: [10.1145/322374.322380](https://doi.org/10.1145/322374.322380). URL: <http://doi.acm.org/10.1145/322374.322380>.
- Chevrou, Florent, Aurélie Hurault, and Philippe Quéinnec (2016). “On the diversity of asynchronous communication”. In: *Formal Aspects Comput.* 28.5, pp. 847–879. DOI: [10.1007/s00165-016-0379-x](https://doi.org/10.1007/s00165-016-0379-x). URL: <https://doi.org/10.1007/s00165-016-0379-x>.

- Courcelle, Bruno (2010). “Special tree-width and the verification of monadic second-order graph properties”. In: *FSTTCS*. Vol. 8. LIPIcs, pp. 13–29.
- Cyriac, Aiswarya, Paul Gastin, and K. Narayan Kumar (2012). “MSO Decidability of Multi-Pushdown Systems via Split-Width”. In: *CONCUR 2012 - Concurrency Theory - 23rd International Conference, CONCUR 2012, Newcastle upon Tyne, UK, September 4-7, 2012. Proceedings*. Ed. by Maciej Koutny and Irek Ulidowski. Vol. 7454. Lecture Notes in Computer Science. Springer, pp. 547–561. DOI: [10.1007/978-3-642-32940-1_38](https://doi.org/10.1007/978-3-642-32940-1_38). URL: https://doi.org/10.1007/978-3-642-32940-1_38.
- Di Giusto, Cinzia, Laetitia Laversa, and Étienne Lozes (2020). “On the k-synchronizability of Systems”. In: *Foundations of Software Science and Computation Structures - 23rd International Conference, FOSSACS 2020, Proceedings*. Ed. by Jean Goubault-Larrecq and Barbara König. Vol. 12077. Lecture Notes in Computer Science. Springer, pp. 157–176. DOI: [10.1007/978-3-030-45231-5_9](https://doi.org/10.1007/978-3-030-45231-5_9). URL: https://doi.org/10.1007/978-3-030-45231-5_9.
- Fidge, Colin J (1988). “Timestamps in Message-Passing Systems That Preserve the Partial Ordering,” in: *Proc. 11th Austral. Comput. Sci. Conf. (ACSC ’88)*, pp. 56–66.
- Genest, Blaise, Dietrich Kuske, and Anca Muscholl (2007). “On Communicating Automata with Bounded Channels”. In: *Fundamenta Informaticae* 80.1-3, pp. 147–167.
- Genest, Blaise, Anca Muscholl, and Dietrich Kuske (2004). “A Kleene Theorem for a Class of Communicating Automata with Effective Algorithms”. In: *Developments in Language Theory, 8th International Conference, DLT 2004, Auckland, New Zealand, December 13-17, 2004, Proceedings*. Ed. by Cristian Calude, Elena Calude, and Michael J. Dinneen. Vol. 3340. Lecture Notes in Computer Science. Springer, pp. 30–48. DOI: [10.1007/978-3-540-30550-7_4](https://doi.org/10.1007/978-3-540-30550-7_4). URL: https://doi.org/10.1007/978-3-540-30550-7_4.
- Giusto, Cinzia Di, Laetitia Laversa, and Étienne Lozes (2021). “Guessing the buffer bound for k-synchronizability”. In: *Implementation and Application of Automata - 25th International Conference, CIAA 2021, Proceedings*. Lecture Notes in Computer Science. To appear. Springer.
- Henriksen, Jesper G. et al. (2005). “A theory of regular MSC languages”. In: *Information and Computation* 202.1, pp. 1–38. ISSN: 0890-5401.
- Kuske, Dietrich and Anca Muscholl (2014). *Communicating automata*.
- Lohrey, Markus and Anca Muscholl (2002). “Bounded MSC Communication”. In: *Foundations of Software Science and Computation Structures, 5th International Conference, FOSSACS 2002. Held as Part of the Joint European Conferences on Theory and Practice of Software, ETAPS 2002 Grenoble, France, April 8-12, 2002, Proceedings*. Ed. by Mogens Nielsen and Uffe Engberg. Vol. 2303. Lecture Notes in Computer Science. Springer, pp. 295–309. DOI: [10.1007/3-540-45931-6_21](https://doi.org/10.1007/3-540-45931-6_21). URL: https://doi.org/10.1007/3-540-45931-6_21.
- (2004). “Bounded MSC communication”. In: *Inf. Comput.* 189.2, pp. 160–181. DOI: [10.1016/j.ic.2003.10.002](https://doi.org/10.1016/j.ic.2003.10.002). URL: <https://doi.org/10.1016/j.ic.2003.10.002>.
- Madhusudan, P. and Gennaro Parlato (2011). “The tree width of auxiliary storage”. In: *Proceedings of the 38th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL 2011, Austin, TX, USA, January 26-28, 2011*. Ed. by Thomas Ball and Mooly Sagiv. ACM, pp. 283–294.
- Mattern, Friedemann (1989). “Virtual time and global states of distributed systems.” In: *Proc. Workshop on Parallel and Distributed Algorithms*, North-Holland / Elsevier, pp. 215–226.
- Schipper, André, Jorge Egli, and Alain Sandoz (1989). “A New Algorithm to Implement Causal Ordering”. In: *Distributed Algorithms, 3rd International Workshop, Nice, France, September 26-28, 1989, Proceedings*. Ed. by Jean-Claude Bermond and Michel Raynal. Vol. 392. Lecture Notes in Computer Science. Springer, pp. 219–232. DOI: [10.1007/3-540-51687-5_45](https://doi.org/10.1007/3-540-51687-5_45). URL: https://doi.org/10.1007/3-540-51687-5_45.